

Corpus Annotation and Data Analysis Equinox
School: Introduction to Statistics

Maik Thalmann
maik.thalmann@gmail.com

September 2022

Contents

1	Prerequisites	5
1.1	R	5
1.2	Statistics	6
1.3	Get in touch	7
	Session Info	7
2	First Session	9

Hey everyone,

if you're looking at this, I assume that you signed up for the CAnDA equinox school in Göttingen in September 2022 and want to attend the Introduction to Statistics course. Here, I will cover some tools and concepts that are going to be instrumental in making sure that you get the most out of this class. I know that it is not ideal to have an introductory class that assumes basic familiarity with some material already, but unfortunately we do not have an entire semester together, but only one short week.

Chapter 1

Prerequisites

Because our schedule for the summer school session on introduction to statistics is quite tight, I will have to presuppose some familiarity with statistics. Here, I will briefly summarize what these requirements are and, if you do not yet feel comfortable with them, give some pointers on where to acquire it.

1.1 R

As you might have guessed, our technical analysis tool will be R. Because an introduction to R would be a class in and of itself, it would be beneficial for all attendees of the class to at least have some basic knowledge. This includes:

- What is an R script?
- How do I create and execute one?
- Which program should I use for my R environment, i.e., for editing R scripts, viewing data and plots, and for running statistical analyses. The popular choice here is undoubtedly RStudio, but you may also use Visual Studio Code and set it up for R if you're more comfortable with that.
- How do I load my data (csv, xlsx, txt file) into R?
- What are factors and integers and how do I switch between them in R?

As an aside, in case I do show code for data manipulation or plots, I will mostly rely on the packages in the so-called tidyverse, a collection of R packages. While I do not consider familiarity with all of these packages

essential, they are important (and often a time saver) independently of this class if you want to use R for your own data analysis or data visualization projects. The packages I will most heavily rely on are dplyr for data wrangling and ggplot2 for visualization purposes.

If you know German (and prefer it over English resources), I have my own website to offer you as a way of (re)gaining familiarity with R. Sessions 1 through 6 should form a quite thorough background (with some skippable material).

Otherwise, I can recommend the relevant chapters in Gries (2013).

1.2 Statistics

As announced in the program for the class, I will, again for reasons of time, have to ask you to know your way around two widely used statistical tests and some basic notions of inferential statistics, detailed below. I am, of course, happy to answer questions during class and the practice sessions, but if you're not as confident with the topics below, I would advise to do some preparatory reading to get the most out of the class (and the one in the second week of the summer school).

- What is the t-test? What does the output of the `t.test` command in R mean – see below?

```
library(tidyverse)
# subset the data
diamond_sub <- diamonds %>%
  filter(color %in% c("E", "J"))
head(diamond_sub)
```

```
# A tibble: 6 x 10
  carat cut      color clarity depth table price      x      y      z
  <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>
1  0.23 Ideal    E     SI2     61.5   55   326   3.95   3.98   2.43
2  0.21 Premium E     SI1     59.8   61   326   3.89   3.84   2.31
3  0.23 Good    E     VS1     56.9   65   327   4.05   4.07   2.31
4  0.31 Good    J     SI2     63.3   58   335   4.34   4.35   2.75
5  0.24 Very Good J     VVS2    62.8   57   336   3.94   3.96   2.48
6  0.22 Fair    E     VS2     65.1   61   337   3.87   3.78   2.49
```

```
t.test(diamond_sub$price ~ diamond_sub$color)
```

```
Welch Two Sample t-test

data:  diamond_sub$price by diamond_sub$color
t = -24.8811, df = 3766.32, p-value < 0.000000000000000222
```

```
alternative hypothesis: true difference in means between group E and group J is not equal to 0
95 percent confidence interval:
 -2424.1311 -2070.0000
sample estimates:
mean in group E mean in group J
 3076.7525      5323.8180
```

- In which scenarios is the *t*-test applicable and when is it not (scale levels, assumptions of the *t*-test, etc.)? What is the effect of setting the `paired` argument in the R command of `t.test` to `TRUE`?
- What is the χ^2 test? Why does it find so much use in corpus linguistics compared to the *t*-test? What does the output below mean?

```
(color_frequencies <- diamond_sub %>%
  group_by(color) %>%
  summarise(frequency = n())
)
```

```
# A tibble: 2 x 2
  color frequency
<ord>      <int>
1 E         9797
2 J         2808
```

```
color_frequencies %>%
  select(frequency) %>%
  chisq.test()
```

Chi-squared test for given probabilities

```
data: .
X-squared = 3875.14, df = 1, p-value < 0.000000000000000222
```

- What are proper and improper interpretations of a *p*-value? What is statistical significance?

1.3 Get in touch

If you have any questions about the information presented here or any other matters related to the class, please do not hesitate to drop me a line via e-mail.

Session Info

```
R version 4.2.1 (2022-06-23)
Platform: x86_64-apple-darwin17.0 (64-bit)
Running under: macOS Big Sur ... 10.16
```

Locale: en_US.UTF-8 / en_US.UTF-8 / en_US.UTF-8 / C / en_US.UTF-8 / en_US.UTF-8

Package version:

forcats_0.5.1	stringr_1.4.0	dplyr_1.0.9	purrr_0.3.4	readr_2.1.2
tibble_3.1.7	ggplot2_3.3.6	tidyverse_1.3.1	styler_1.7.0	tidyselect_1.1
rematch2_2.1.2	haven_2.5.0	colorspace_2.0-3	vctrs_0.4.1	generics_0.1.3
yaml_2.3.5	utf8_1.2.2	rlang_1.0.3	R.oo_1.25.0	pillar_1.7.0
glue_1.6.2	withr_2.5.0	DBI_1.1.3	dbplyr_2.2.1	modelr_0.1.8
R.cache_0.15.0	lifecycle_1.0.1	munsell_0.5.0	gtable_0.3.0	cellranger_1.1
rvest_1.0.2	codetools_0.2-18	evaluate_0.15	knitr_1.39	tzdb_0.3.0
fansi_1.0.3	broom_1.0.0	backports_1.4.1	scales_1.2.0	jsonlite_1.8.0
hms_1.1.1	digest_0.6.29	stringi_1.7.6	bookdown_0.27	grid_4.2.1
tools_4.2.1	magrittr_2.0.3	crayon_1.5.1	pkgconfig_2.0.3	ellipsis_0.3.2
reprex_2.0.1	lubridate_1.8.0	assertthat_0.2.1	rmarkdown_2.14	httr_1.4.3
R6_2.5.1	compiler_4.2.1			

Chapter 2

First Session

(to be populated)

Gries, Stefan T. 2013. *Statistics for linguistics with R: A practical introduction*. Berlin: de Gruyter.